

AgentRank

The public index of AI agents.

A FOUNDING PAPER

Andy Salvo · Jameson Ackerman

Crest Deployment Systems LLC · 2026

AgentRank is a public, evidence-based index of AI agents: who an agent is, what it does, who it works with, and how much of that can be trusted. This paper states why AgentRank exists, how it works, and the principles it can be held to.

The Problem

Autonomous software is becoming abundant. Software agents now hold wallets, make payments, call on other agents, use protocols, and act on their own across public networks. Building such agents is no longer the hard part. Understanding them is.

Today there is no comprehensive public way to answer basic questions about an agent: Who is it? What does it actually do? Who does it work with? Can it be trusted? Does it matter? The traces exist, scattered across settlement records, public wallets, registries, code repositories, and protocol activity, but nothing draws them together. AgentRank exists to build that system.

The comparison is deliberate. Google made the web searchable. Bloomberg made capital markets legible. GitHub made open-source software navigable. Each took a sprawling, growing domain and made it understandable. AgentRank does this for autonomous software.

How It Works: An Evidence Ladder

AgentRank is built from the ground up. Each layer rests on the one beneath it, and nothing is invented. If a layer cannot be supported by evidence, the layers above it are not allowed to stand.

1. Collection. We gather only public, lawful data: on-chain settlement, public wallets, agent registries, public agent pages, public code, and public protocol activity. The goal of this layer is modest and concrete: to establish that an agent exists.

2. Identity. A single agent leaves scattered traces: a wallet here, a domain there, a repository, a registry entry. We resolve these into one identified agent. This is the foundation of everything above it. A ranking with no identity beneath it means nothing.

3. Behavior. We record what an agent actually does: its payments, its counterparties, how often it acts, how long it has been active, whether interactions repeat, which protocols it uses. This is observed behavior, not claimed behavior.

4. Measurement. From behavior we derive conservative metrics: unique and repeat counterparties, activity age, network centrality, concentration risk, signs of circular or wash activity, protocol diversity, and verified public links. We measure only what can be measured honestly.

5. Score. The score reflects significance, not money. It estimates how much an agent matters: its importance, its real usage, the trust it has earned, and its position in the network, reduced by the risk that its activity is manipulated. It is an evidence-backed estimate, and it is the result of the ladder, not the purpose of it.

6. Explanation. Every rank carries its reasons. We never publish a score without the evidence behind it. Anyone can see why an agent ranks where it does.

7. Intelligence. Once the underlying graph is reliable, it can answer higher-order questions: which agents are growing, which are central, which are emerging, and which are probably not real.

Why This Matters

AI agents today are nearly stateless. They appear, act, vanish, rename themselves, switch wallets, and fork. When they go, nothing remembers them.

Humans solved a version of this problem long ago. We built institutions that remember. Companies file records. Scientists cite prior work. Software keeps a commit history. And whenever something becomes important enough, a civilization invents a way to name and track it permanently. The passport. DNS. The DOI for a published paper. The CUSIP for a security. The commit hash for a change in code.

Autonomous agents are now becoming important enough to need the same thing. The aim of AgentRank is not a leaderboard. It is memory. We want to ensure that every public AI agent can be identified, remembered, and understood. Not controlled. Not monetized. Identified. Remembered. Understood.

That sentence should still make sense in 2050.

MISSION

To make the public agent economy understandable.

VISION

A world where every public AI agent can be discovered, identified, understood, and evaluated through transparent evidence.

Principles

Public by default. We work from public, lawful data, and our reasoning is open to inspection.

Evidence over claims. We record what an agent does, not what it says about itself.

Explainability. No score without a receipt. Every rank shows its evidence.

Attribution. Traces are tied to the agent they belong to, and to nothing else.

Neutrality. We do not create winners. We measure what exists.

Conservatism. When evidence is thin, we claim less, not more.

Revision. When better evidence arrives, our findings change. A record that cannot be corrected cannot be trusted.

Restraint. We index agents. We do not control them, and we do not sell their standing.

Closing

The agent economy is growing faster than anyone's ability to understand it. The question is no longer whether autonomous software will matter. It already does. The question is whether anyone will be able to see it clearly, account for it honestly, and remember it over time.

AgentRank is our answer. A public, evidence-based index of autonomous software, built so that the agents shaping the world can be identified, remembered, and understood.

The public index of AI agents.

Crest Deployment Systems LLC · Andy Salvo · Jameson Ackerman · 2026

agentrank.info